



NLP Tweet Analysis

Will Byrd June 2024

Introduction

In this NLP project, we will explore preprocessing of text data and sophisticated White Box modeling approaches to determine the **sentiment** of tweets. We can tell based on a quick review of the dataset that these are tweets made from attendees to South by Southwest or SXSW. For context, SWSW is conference that celebrates technology and arts. From a business perspective, we have been tasked with determining which brands are were rated most favorably so SXSW will know who to allocate resources to vendors next year.

In this notebook, we will use preprocessing tools such as:

- regex - text data cleaning tool
- stemming - stripping affixes from words-leaving base forms
- lemmatization - ensuring the output word is a normalized version of the word
- label encoding - converting categorical variables into numerical format
- imputing - correcting NaNs
- tokenization - splitting text into smaller units such as words or bigrams/trigrams
- vectorization - converting text into numerical representations for modeling

Data

Our analysis will be performed on a csv file from CrowdFlower via data.world. This file contains over 9000 tweets that have been rated by humans to be positive, negative, or nuetral. We also have insight into which brand or product is being targeted by each tweet ('emotion_in_tweet_is_directed_at') that can be analyzed during our EDA. Ultimately, we will end with 3 features that we will use to predict sentiment:

- pos_tagged_text - strings of tokenized text with parts of speech labelled to each word
- bigrams - a column containing all bigrams (pairs of words)
- trigrams - a column containing all trigrams (words occuring in groups of 3)

Goals

Our goal will be to build a model that can most accurately predict **Sentiment** of these tweets as it relates to specific brands. Sentiment will be encoded to have a score between 1-3:

- 1 negative
- 2 indifferent
- 3 positive

Since we already have a robust dataset with **Sentiment** (originally labelled as 'is_there_an_emotion_directed_at_a_brand_or_product'), we will build a supervised learning model using this df for training and testing.

Overview



Let's take a look at the data to better understand what we need to do to the text data to analyze and model it. Imputing the brand column that is labelled 'emotion_in_tweet_is_directed_at' will also be important to improve the sample size and power of our results. We will build a variety of models:

- Logistic Regression Model
- Decision Tree Classifier
- K-Nearest Neighbors
- Gradient Boosting

After building these models and finetuning results, Stacking and Voting Ensemble methods are then used to further improve our models. This is a costly process and takes our machine quite some time to run, but it ensures our results are optimal.


Importing all of the necessary libraries.

In [1]: ▶

```
1 import pandas as pd
2 import re
3 import numpy as np
4 import nltk
5 import seaborn as sns
6 from sklearn.preprocessing import LabelEncoder, OneHotEncoder
7 from sklearn.model_selection import train_test_split, GridSearchCV
8 from sklearn.tree import DecisionTreeClassifier
9 from sklearn.feature_extraction.text import TfidfVectorizer, CountVec
10 from sklearn.linear_model import LogisticRegression
11 from sklearn.metrics import accuracy_score, classification_report, co
12 from sklearn.compose import ColumnTransformer
13 from sklearn.neighbors import KNeighborsClassifier
14 from sklearn.ensemble import StackingClassifier, GradientBoostingClas
15 from sklearn.pipeline import Pipeline
16 from nltk.corpus import stopwords
17 from nltk.stem import WordNetLemmatizer
18 import matplotlib.pyplot as plt
19 %matplotlib inline
20 from nltk.tokenize import word_tokenize
21 from nltk.corpus import stopwords
22 from nltk.util import ngrams
23 from collections import Counter
24 nltk.download('punkt')
25 nltk.download('averaged_perceptron_tagger')
26 nltk.download('punkt', quiet=True)
27 np.random.seed(0)
28 from mlxtend.plotting import plot_decision_regions
29
```

```
[nltk_data] Downloading package punkt to
[nltk_data]   C:\Users\byrdw\AppData\Roaming\nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data]   C:\Users\byrdw\AppData\Roaming\nltk_data...
[nltk_data]   Package averaged_perceptron_tagger is already up-to-
[nltk_data]   date!
```

Reading in the csv tweets.csv to inspect specific rows and columns.

In [3]: 

```
1 df = pd.read_csv('tweets.csv')
2 df
```



Out[3]:

	tweet_text	emotion_in_tweet_is_
0	.@wesley83 I have a 3G iPhone. After 3 hrs twe...	
1	@jessedee Know about @fludapp ? Awesome iPad/i...	iPad or
2	@swonderlin Can not wait for #iPad 2 also. The...	
3	@sxsw I hope this year's festival isn't as cra...	iPad or
4	@sxtxstate great stuff on Fri #SXSW: Marissa M...	
...	...	
8716	Ipad everywhere. #SXSW {link}	
8717	Wave, buzz... RT @mention We interrupt your re...	
8718	Google's Zeiger, a physician never reported po...	
8719	Some Verizon iPhone customers complained their...	

Let's rename the column that contains all of the brands/products, since we will be working with that column throughout this notebook.

In [4]: 

```
1 df.rename(columns={'emotion_in_tweet_is_directed_at': 'brand'}, inplace=True)
```

Let's take a quick look at the info of this df. We can see lots of missing values in the **'brand'** column. We can probably impute most of these from the **'tweet_text'** column.

In [5]: 

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8721 entries, 0 to 8720
Data columns (total 3 columns):
 #   Column                                Non-Null Count
Dtype
---  ---
0    tweet_text                          8720 non-null
object
1    brand                               3169 non-null
object
2    is_there_an_emotion_directed_at_a_brand_or_product  8721 non-null
object
dtypes: object(3)
memory usage: 204.5+ KB
```



Exploratory Data Analysis

Now that the dataset is loaded, we can begin our EDA. First thing to address is the imputation of values for our **'brand'** column. Let's look at all unique values.

Looking at unique values in column 'emotion_in_tweet_is_directed_at' for imputation.

```
In [6]: 1 df['brand'].unique()
```

```
Out[6]: array(['iPhone', 'iPad or iPhone App', 'iPad', 'Google', nan, 'Android',  
              'Apple', 'Android App', 'Other Google product or service',  
              'Other Apple product or service'], dtype=object)
```

Found some NaN values in our 'brand' column and want to impute some additional values where possible to reduce the amount of NaN values. Looks like iPhone, iPad or iPad App, iPad, Google, Android, Apple, Android App, Other Google product or service, and Other Apple product or service are all of our values currently.

Since we are more concerned with major brands and products, let's consolidate this list to:

- iPhone
- iPad
- Apple
- Google
- Android

In [7]:

```

1 # Define the list of words we want to check for in the 'tweet_text'
2 words_to_check = ['iPhone', 'Apple', 'Google', 'iPad', 'Android']
3
4 # Replace NaN values in 'tweet_text' column with an empty string
5 df['tweet_text'] = df['tweet_text'].fillna('')
6
7 # Filter the DataFrame to include only rows where 'brand' is NaN
8 filtered_df = df[df['brand'].isna()]
9
10 # Loop over each word to check for in the 'tweet_text' column
11 for word in words_to_check:
12     # Use boolean indexing to find rows where 'tweet_text' contains the word
13     rows_with_word = filtered_df[filtered_df['tweet_text'].str.contains(word)]
14
15     # Update the value of 'brand' for the matching rows
16     df.loc[rows_with_word.index, 'brand'] = word
17
18 # Print the updated DataFrame
19 print(df[['tweet_text', 'brand']])

```

	tweet_text	brand
0	.@wesley83 I have a 3G iPhone. After 3 hrs twe...	iPhone
1	@jessedee Know about @fludapp ? Awesome iPad/i...	iPad or iPhone
2	@swonderlin Can not wait for #iPad 2 also. The...	iPad
3	@sxsxw I hope this year's festival isn't as cra...	iPad or iPhone
4	@sxtxstate great stuff on Fri #SXSW: Marissa M...	Google
...	...	
8716	Ipad everywhere. #SXSW {link}	iPad
8717	Wave, buzz... RT @mention We interrupt your re...	Google
8718	Google's Zeiger, a physician never reported po...	Google
8719	Some Verizon iPhone customers complained their...	iPhone
8720	Google	


[8721 rows x 2 columns]

In [8]:  1 df.info() *# checking the info again to see how many values we were d*

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8721 entries, 0 to 8720
Data columns (total 3 columns):
#   Column                                Non-Null Count
Dtype
---  ---
0    tweet_text                          8721 non-null
object
1    brand                              7984 non-null
object
2    is_there_an_emotion_directed_at_a_brand_or_product  8721 non-null
object
dtypes: object(3)
memory usage: 204.5+ KB
```

Pretty solid improvement from 3169 values to 7984 values now. Let's use mapping to consolidate these down to our targeted list we mentioned earlier:

- iPhone
- iPad
- Apple
- Google
- Android

In [9]:  1 *# taking a quick look at unique value counts and we can see the value*
2 df['brand'].value_counts()

```
Out[9]: iPad                2389
Google                2057
Apple                1299
iPhone                961
iPad or iPhone App    451
Android              433
Other Google product or service  282
Android App           78
Other Apple product or service  34
Name: brand, dtype: int64
```

```

In [10]: 1 # Define the list of categories we keep
2 categories_to_keep = ['iPhone', 'iPad', 'Apple', 'Google', 'Android']
3
4 # Define a mapping of possible values to the categories we want to keep
5 mapping = {
6     'iphone': 'iPhone',
7     'ipad': 'iPad',
8     'apple': 'Apple',
9     'google': 'Google',
10    'android': 'Android',
11    'android app': 'Android',
12    'ipad or iphone app': 'Apple',
13    'other apple product or service': 'Apple',
14    'other google product or service': 'Google'
15 }
16
17 # Convert all values to lowercase for case-insensitive matching
18 df['brand'] = df['brand'].str.lower()
19
20 # Map the values to the desired categories using the mapping defined
21 df['brand'] = df['brand'].apply(lambda x: mapping.get(x, x))
22
23 # Replace any remaining empty strings or NaNs with NaN
24 df['brand'].replace('', pd.NA, inplace=True)
25
26 # Print the updated DataFrame
27 print(df['brand'].value_counts())
28

```

```

iPad      2389
Google    2339
Apple     1784
iPhone     961
Android    511
Name: brand, dtype: int64

```

To perform analysis and build models, we will need to standardize all of our text data. Standardizing the data includes:

- making everything lowercase
- removing nonword characters and symbols
- stripping whitespaces
- removing stop words
- Lemmatization
- Stemming

Now, let's define a function that will use regex, lemmatization and stemming to clean our data. For context, lemmatization and stemming is a process that essentially reduces words down to their bases. Running and runs become run. Partying and parties become party, etc.



```
In [11]: ▶ 1 # Text cleaning using regex
2 def clean_tweet(tweet):
3     tweet = re.sub(r'http\S+|www\S+|https\S+', '', tweet, flags=re.M)
4     tweet = re.sub(r'@\w+|\#', '', tweet) ## removing @, #
5     tweet = tweet.lower() # make all text lower case
6     tweet = re.sub(r'\W', ' ', tweet) # replace non word characters with space
7     tweet = re.sub(r'\s+', ' ', tweet) # replace multiple spaces with single space
8     tweet = tweet.strip() # removes leading whitespaces
9     stop_words = set(stopwords.words('english')) # removing stop words
10    lemmatizer = WordNetLemmatizer() # lemmatizing words, turning running into run
11    words = tweet.split() # creating a list of words
12    cleaned_words = [lemmatizer.lemmatize(word) for word in words if word not in stop_words]
13    return ' '.join(cleaned_words) # joining clean and lemmatized words
```

Now we will need to turn our columns into strings for analysis/modeling.

```
In [12]: ▶ 1 df['tweet_text'] = df['tweet_text'].astype(str)
```

```
In [13]: ▶ 1 df['cleaned_tweet'] = df['tweet_text'].apply(clean_tweet)
```

Let's take a look at our work in the new column we have created called **'cleaned_tweet'**. This is a good chance to take a look at what the stemming and lemmatization has done to our text data.

In [14]:  1 df.head()

Out[14]:

	tweet_text	brand	is_there_an_emotion_directed_at_a_brand_or_product	cleaned_tweet
0	.@wesley83 I have a 3G iPhone. After 3 hrs twe...	iPhone	Negative emotion	3g iphone 3 hr tweeting rise_austin dead need ...
1	@jessedee Know about @fludapp ? Awesome iPad/i...	Apple	Positive emotion	know awesome ipad iphone app likely appreciate...
2	@swonderlin Can not wait for #iPad 2 also. The...	iPad	Positive emotion	wait ipad 2 also sale sxsw
3	@sxsw I hope this year's festival isn't as cra...	Apple	Negative emotion	hope year festival crashy year iphone app sxsw
4	@sxtxstate great stuff on Fri #SXSW: Marissa M...	Google	Positive emotion	great stuff fri sxsw marissa mayer google tim ...

Tweets are cleaned and we can move on to the next step-encoding.


Label Encoding values in the 'is_there_an_emotion_directed_at_a_brand_or_product' column to make a new column, 'sentiment' is appropriate since the values are ordinal.

0=Undefined

1=Negative

2=Indifferent

3=Positive

In [15]:  1 label_encoder = LabelEncoder()
2 df['sentiment'] = label_encoder.fit_transform(df['is_there_an_emotion...

Looks like we have some more rows we can drop. If we can't decipher sentiment (denoted as a value of 0), we can't use these in our modeling. Goodbye!

```
In [16]: 1 df['sentiment']
        2 df[df['sentiment']==0]
```

Out[16]:

	tweet_text	brand	is_there_an_emotion_directed_at_a_brand_or_product	c
88	Thanks to @mention for publishing the news of ...	NaN	I can't tell	tha
100	💎💎💎@mention "Apple has opened a pop- up st...	iPad	I can't tell	ope au
228	Just what America needs. RT @mention Google to...	Google	I can't tell	a
330	The queue at the Apple Store in Austin is FOUR...	Apple	I can't tell	st

```
In [17]: 1 # Drop rows where 'sentiment' is equal to 0 in the original DataFrame
        2 df.drop(df[df['sentiment'] == 0].index, inplace=True)
        3
```

Let's take another look to make sure all NaN values have been addressed.

```
In [18]: 1 # quick way to view all NaN values in a specific column
        2 df['brand'].isna().sum()
```

Out[18]: 732

But I'm still seeing some NaN values. Let's finish cleaning this up.

```
In [19]: 1 df.dropna(subset=['brand'], inplace=True)
```

```
In [20]: 1 df.isnull().sum()
```

```
Out[20]: tweet_text      0
        brand            0
        is_there_an_emotion_directed_at_a_brand_or_product  0
        cleaned_tweet    0
        sentiment        0
        dtype: int64
```

```
In [21]: 1 # we need to reset our index for matching in our modelling.
        2 df.reset_index(drop=True, inplace=True)
```

For our **'brand'** column we will need to use **OneHotEncoder** to address the categorical variable in that column. Notice since there is no ordinal relationship between the values in this column, we have to use OneHotEncoding vs label encoding.



```
In [22]: ▶ 1 # Use OneHotEncoder from scikit-Learn
2 encoder = OneHotEncoder(categories=[[ 'iPhone', 'iPad', 'Apple', 'Goog
3 encoded_data = encoder.fit_transform(df[['brand']])
4
5 # Create a DataFrame with the encoded data
6 encoded_df = pd.DataFrame(encoded_data, columns=encoder.get_feature_n
7
8 # Concatenate the encoded DataFrame with the original DataFrame
9 df = pd.concat([df, encoded_df], axis=1)
10
11 # Drop the original 'brand' column if no longer needed
12 #df.drop(columns=['brand'], inplace=True)
13
14 # Print the updated DataFrame
15 df.head()
16
```

C:\Users\byrdw\anaconda3\envs\learn-env\lib\site-packages\sklearn\preprocessing_encoders.py:975: FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its default value.

```
warnings.warn(
```

Out[22]:

	tweet_text	brand	is_there_an_emotion_directed_at_a_brand_or_product	cleaned_tweet
0	.@wesley83 I have a 3G iPhone. After 3 hrs twe...	iPhone	Negative emotion	3g iphone 3 hr tweeting rise_austin dead need ...
1	@jessedee Know about @fludapp ? Awesome iPad/i...	Apple	Positive emotion	know awesome ipad iphone app likely appreciate...
2	@swonderlin Can not wait for #iPad 2 also. The...	iPad	Positive emotion	wait ipad 2 also sale sxsw
3	@sxsw I hope this year's festival isn't as cra...	Apple	Negative emotion	hope year festival crashy year iphone app sxsw
4	@sxtxstate great stuff on Fri #SXSW: Marissa M...	Google	Positive emotion	great stuff fri sxsw marissa mayer google tim ...

One last review for NaN values



```
In [23]: 1 nan_values = df.isna().any()
          2
          3 # Print the columns with NaN values, if any
          4 print(nan_values)
```

tweet_text	False
brand	False
is_there_an_emotion_directed_at_a_brand_or_product	False
cleaned_tweet	False
sentiment	False
brand_iPhone	False
brand_iPad	False
brand_Apple	False
brand_Google	False
brand_Android	False
dtype: bool	

Double check!


```
In [24]: 1 df.isnull().sum()
```

```
Out[24]: tweet_text      0
          brand          0
          is_there_an_emotion_directed_at_a_brand_or_product  0
          cleaned_tweet  0
          sentiment      0
          brand_iPhone   0
          brand_iPad     0
          brand_Apple    0
          brand_Google   0
          brand_Android  0
          dtype: int64
```

Great work! All NaN values are gone and we have imputed 4000 values!

Now it's time to Tokenize our text. Tokenizing is an important preprocessing step as it allows us to further analyze the text. We can create bigrams, trigrams, and determine feature importance.

tokenizing 'cleaned_tweet' to 'tokenized_tweets'

In [25]: 

```
1 df['tokenized_tweets'] = df['cleaned_tweet'].apply(word_tokenize)
2 df.head()
```



Out[25]:

	tweet_text	brand	is_there_an_emotion_directed_at_a_brand_or_product	cleaned_tweet
0	.@wesley83 I have a 3G iPhone. After 3 hrs twe...	iPhone	Negative emotion	3g iphone 3 hr tweeting rise_austin dead need ...
1	@jessedee Know about @fludapp ? Awesome iPad/i...	Apple	Positive emotion	know awesome ipad iphone app likely appreciate...
2	@swonderlin Can not wait for #iPad 2 also. The...	iPad	Positive emotion	wait ipad 2 also sale sxsw
3	@sxsw I hope this year's festival isn't as cra...	Apple	Negative emotion	hope year festival crashy year iphone app sxsw
4	@sxtxstate great stuff on Fri #SXSW: Marissa M...	Google	Positive emotion	great stuff fri sxsw marissa mayer google tim ...



Pretty simple to tokenize the text. Now that the individual words are tokenized, let's create a bag of words. This will allow us to check frequency of words.

```
In [26]: ▶ 1 def count_vectorize(tokenized_list): # Define the function with a l
2 corpus = {} # Initialize an empty dictionary to store word counts
3
4     for tokenized in tokenized_list: # Iterate over each tokenized tweet
5         for word in tokenized: # Iterate over each word in the tokenized tweet
6             if word in corpus: # If the word is already in the dictionary
7                 corpus[word] += 1 # Increment its count by 1
8             else: # If the word is not in the dictionary
9                 corpus[word] = 1 # Add it to the dictionary with a count of 1
10
11     return corpus # Return the dictionary containing word counts
12
13 # Vectorize the tokenized tweets
14 text_vectorized = count_vectorize(df['tokenized_tweets']) # Call the function
15 text_vectorized # Output the resulting dictionary
```

```
Out[26]: {'3g': 29,
'iphone': 1492,
'3': 143,
'hr': 5,
'tweeting': 28,
'rise_austin': 2,
'dead': 7,
'need': 218,
'upgrade': 13,
'plugin': 4,
'station': 9,
'sxsw': 8324,
'know': 181,
'awesome': 115,
'ipad': 2397,
'app': 751,
'likely': 11,
'appreciate': 4,
'design': 130,
...}
```

checking frequency of words

```
In [27]: 1 freq_df = pd.DataFrame(list(text_vectorized.items()), columns=['Word', 'Frequency'])
        2 freq_df
```



Out[27]:

	Word	Frequency
0	3g	29
1	iphone	1492
2	3	143
3	hr	5
4	tweeting	28
...
8027	complained	1
8028	yorkers	1
8029	٣	1
8030	u	1
8031	٥	1


8032 rows × 2 columns

Let's visualize the tope 20 most frequent words

In [28]:



```
1 # Convert the word frequency dictionary to a DataFrame
2 freq_df = pd.DataFrame(list(text_vectorized.items()), columns=['Word', 'Frequency'])
3
4 # Sort the DataFrame by frequency
5 freq_df = freq_df.sort_values(by='Frequency', ascending=False)
6
7 # Plotting
8 plt.figure(figsize=(10, 6))
9 sns.barplot(x='Frequency', y='Word', data=freq_df.head(20), palette='magma')
10 plt.title('Top 20 Most Frequent Words')
11 plt.xlabel('Frequency')
12 plt.ylabel('Words')
13 plt.show()
```



Now we can take a look at this a different way- let's look at the normalized frequency of each word



In [29]:



```

1 # Calculate the total number of tokens
2 total_word_count = freq_df['Frequency'].sum()
3
4 # Print the total number of tokens
5 print("Total word count:", total_word_count)
6
7 # Extract the top 50 most common words
8 df_top_50 = freq_df.head(50)
9
10 # Print the top 50 words and their normalized frequencies
11 print(f'{"Word":10} Normalized Frequency')
12 for word in df_top_50.iterrows():
13     normalized_frequency = word[1]['Frequency'] / total_word_count
14     print(f'{"Word":10} {normalized_frequency:.6f}')
15

```

Total word count: 92923

Word	Normalized Frequency
sxsw	0.089580
link	0.036676
rt	0.028389
google	0.026915
ipad	0.025796
apple	0.023460
iphone	0.016056
quot	0.015906
store	0.015238
2	0.012483
new	0.010977
austin	0.009147
app	0.008082
amp	0.007695
circle	0.006898
launch	0.006887
social	0.006457
...	...

This tells us that SXSW is the most commonly occurring word. Another interesting thing to notice is the order in which our main brands/product appear. This is the percentage of words that in the corpus that are the specific word. So out of nearly 100,000 words-8.9% of them are sxsw.

- google 2.6915%
- ipad 2.5796%
- apple 2.3460%
- iphone 1.6056%
- android 0.6037%

Let's get a total count for how many tweets talk about each specific brand. Remember there are 7984 tweets with a specific brand focus.



```
In [30]: 1 # Calculating the total count of tweets directed at Google
2 google_count = df[df['brand'] == 'Google']['brand'].count()
3 google_count
```

Out[30]: 2289

```
In [31]: 1 # Calculating percentage of tweets that are about Google
2 google_count = df['brand'].value_counts().get('Google', 0)
3 total_count = df['brand'].count()
4 google_percent = google_count / total_count
5 google_percent
6
```

Out[31]: 0.2920387854044399

```
In [32]: 1 # Calculating the total count of tweets directed at Apple
2 apple_count = df[df['brand'] == 'Apple']['brand'].count()
3 apple_count
```

Out[32]: 1764

```
In [33]: 1 # Calculating percentage of tweets that are about Apple
2 apple_count = df['brand'].value_counts().get('Apple', 0)
3 total_count = df['brand'].count()
4 apple_percent = apple_count / total_count
5 apple_percent
6
```

Out[33]: 0.22505741260525644

```
In [34]: 1 # Calculating the total count of tweets directed at iPad
2 iPad_count = df[df['brand'] == 'iPad']['brand'].count()
3 iPad_count
```

Out[34]: 2346

```
In [35]: 1 # Calculating percentage of tweets that are about iPad
2 iPad_count = df['brand'].value_counts().get('iPad', 0)
3 total_count = df['brand'].count()
4 iPad_percent = iPad_count / total_count
5 iPad_percent
6
```

Out[35]: 0.29931104873692266

```
In [36]: 1 # Calculating the total count of tweets directed at iPhone
2 iPhone_count = df[df['brand'] == 'iPhone']['brand'].count()
3 iPhone_count
```

Out[36]: 934



```
In [37]: ▶ 1 # Calculating percentage of tweets that are about iPhone
2 iPhone_count = df['brand'].value_counts().get('iPhone', 0)
3 total_count = df['brand'].count()
4 iPhone_percent = iPhone_count / total_count
5 iPhone_percent
6
```

Out[37]: 0.1191630517989283

```
In [38]: ▶ 1 # Calculating the total count of tweets directed at Android
2 Android_count = df[df['brand'] == 'Android']['brand'].count()
3 Android_count
```

Out[38]: 505

```
In [39]: ▶ 1 # Calculating percentage of tweets that are about Android
2 Android_count = df['brand'].value_counts().get('Android', 0)
3 total_count = df['brand'].count()
4 Android_percent = Android_count / total_count
5 Android_percent
6
```

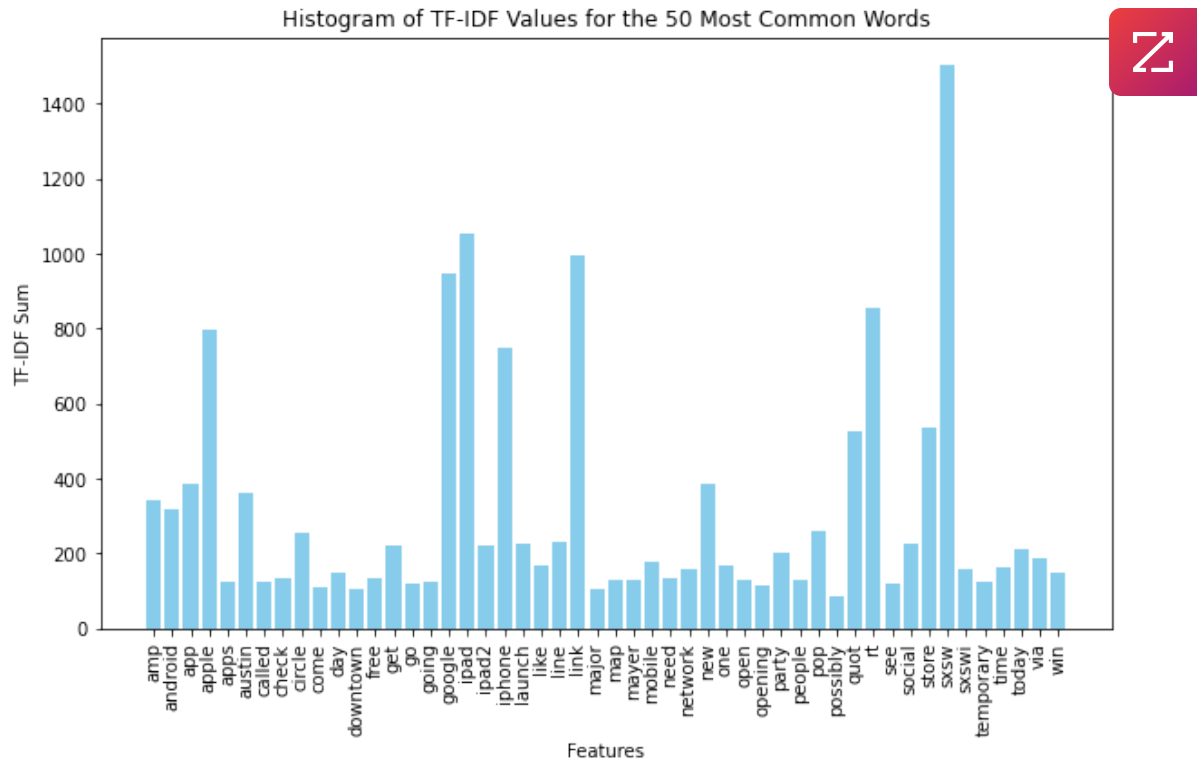
Out[39]: 0.06442970145445266

Now let's take a look at the histogram of TF-IDF Values for the 50 most common words



In [40]: ▶

```
1 # Ensure 'cleaned_tweet' column is of string type
2 df['cleaned_tweet'] = df['cleaned_tweet'].astype(str)
3
4 # Apply the clean_tweet function to each entry in the 'cleaned_tweet'
5 df['cleaned_text'] = df['cleaned_tweet'].apply(clean_tweet)
6
7 # Convert the cleaned text column to a list
8 text_data = df['cleaned_text'].tolist()
9
10 # Initialize a CountVectorizer with a maximum of 50 features
11 count_vectorizer = CountVectorizer(max_features=50)
12
13 # Fit and transform the text data to a document-term matrix
14 X_counts = count_vectorizer.fit_transform(text_data)
15
16 # Get the most common words (features) from the count vectorizer
17 common_words = count_vectorizer.get_feature_names_out()
18
19 # Initialize a TfidfVectorizer with the common words as vocabulary
20 vectorizer_tfidf = TfidfVectorizer(vocabulary=common_words)
21
22 # Fit and transform the text data to a TF-IDF matrix
23 X_tfidf = vectorizer_tfidf.fit_transform(text_data)
24
25 # Convert the TF-IDF matrix to an array
26 tfidf_matrix = X_tfidf.toarray()
27
28 # Get the feature names (words) from the TF-IDF vectorizer
29 feature_names = vectorizer_tfidf.get_feature_names_out()
30
31 # Sum the TF-IDF values for each feature across all documents
32 tfidf_sums = np.sum(tfidf_matrix, axis=0)
33
34 # Plot a histogram of the TF-IDF values for the 50 most common words
35 plt.figure(figsize=(10, 6))
36 plt.bar(feature_names, tfidf_sums, color='skyblue')
37 plt.xlabel('Features')
38 plt.ylabel('TF-IDF Sum')
39 plt.title('Histogram of TF-IDF Values for the 50 Most Common Words')
40 plt.xticks(rotation=90)
41 plt.show()
42
```



Let's add Part of Speech Tagging here as well. This is another level preprocessing and can improve our model.

```
In [41]: 1 def pos_tagging(text):
2         # Tokenize the input text into individual words
3         words = nltk.word_tokenize(text)
4
5         # Assign part-of-speech tags to each word
6         pos_tags = nltk.pos_tag(words)
7
8         # Format each word and its POS tag as 'word_tag' and join them in
9         return ' '.join([f"{word}_{tag}" for word, tag in pos_tags])
10
11 # Apply the pos_tagging function to each entry in the 'cleaned_text'
12 # and store the result in a new column 'pos_tagged_text'
13 df['pos_tagged_text'] = df['cleaned_text'].apply(pos_tagging)
14
```

In [42]:

```
1 # Inspecting our df after adding some columns
2 df.head()
```



Out[42]:

	tweet_text	brand	is_there_an_emotion_directed_at_a_brand_or_product	cleaned_tweet
0	.@wesley83 I have a 3G iPhone. After 3 hrs twe...	iPhone	Negative emotion	3g iphone 3 hr tweeting rise_austin dead need ...
1	@jessedee Know about @fludapp ? Awesome iPad/i...	Apple	Positive emotion	know awesome ipad iphone app likely appreciate...
2	@swonderlin Can not wait for #iPad 2 also. The...	iPad	Positive emotion	wait ipad 2 also sale sxsw
3	@sxsw I hope this year's festival isn't as cra...	Apple	Negative emotion	hope year festival crashy year iphone app sxsw
4	@sxtxstate great stuff on Fri #SXSW: Marissa M...	Google	Positive emotion	great stuff fri sxsw marissa mayer google tim ...

Now that we have created our bag of words, let's create bigrams and trigrams. These can be important in fully understanding meaning in text. For example-which iPad is everyone talking about?



```
In [43]: ▶ 1 def generate_ngrams(text, n):
2         tokens = word_tokenize(text)
3         n_grams = list(ngrams(tokens, n))
4         return [' '.join(gram) for gram in n_grams]
5
6         # Apply the function to generate bigrams and trigrams
7 df['bigrams'] = df['cleaned_text'].apply(lambda x: generate_ngrams(x,
8 df['trigrams'] = df['cleaned_text'].apply(lambda x: generate_ngrams(x
9
10 # Display the DataFrame with bigrams and trigrams
11 df['bigrams']
```

```
Out[43]: 0      [3g iphone, iphone 3, 3 hr, hr tweeting, tweet...
1      [know awesome, awesome ipad, ipad iphone, ipho...
2      [wait ipad, ipad 2, 2 also, also sale, sale sxsw]
3      [hope year, year festival, festival crashy, cr...
4      [great stuff, stuff fri, fri sxsw, sxsw mariss...
          ...
7833      [ipad everywhere, everywhere sxsw, sxsw link]
7834      [wave buzz, buzz rt, rt interrupt, interrupt r...
7835      [google zeiger, zeiger physician, physician ne...
7836      [verizon iphone, iphone customer, customer com...
7837      [? _, _ u, u й, й _, _ _, _ rt, rt google, goo...
Name: bigrams, Length: 7838, dtype: object
```


Let's create a graph to view the top 50 most occurring bigrams. This can give insight into which products people are talking about.

In [44]:

```

1 # Flatten the list of bigrams
2 all_bigrams = [bigram for sublist in df['bigrams'] for bigram in sublist]
3
4 # Count the frequency of each bigram
5 bigram_freq = Counter(all_bigrams)
6
7 # Get the top 50 most common bigrams
8 top_50_bigrams = bigram_freq.most_common(50)
9
10 # Prepare data for the histogram
11 bigrams, counts = zip(*top_50_bigrams)
12 bigram_labels = [' '.join(bigram) for bigram in bigrams]
13
14 # Plot the histogram
15 plt.figure(figsize=(12, 8))
16 plt.barh(bigram_labels, counts, color='skyblue')
17 plt.xlabel('Frequency')
18 plt.ylabel('Bigrams')
19 plt.title('Top 50 Most Common Bigrams')
20 plt.gca().invert_yaxis() # Invert y-axis to have the highest frequency at the top
21 plt.show()

```



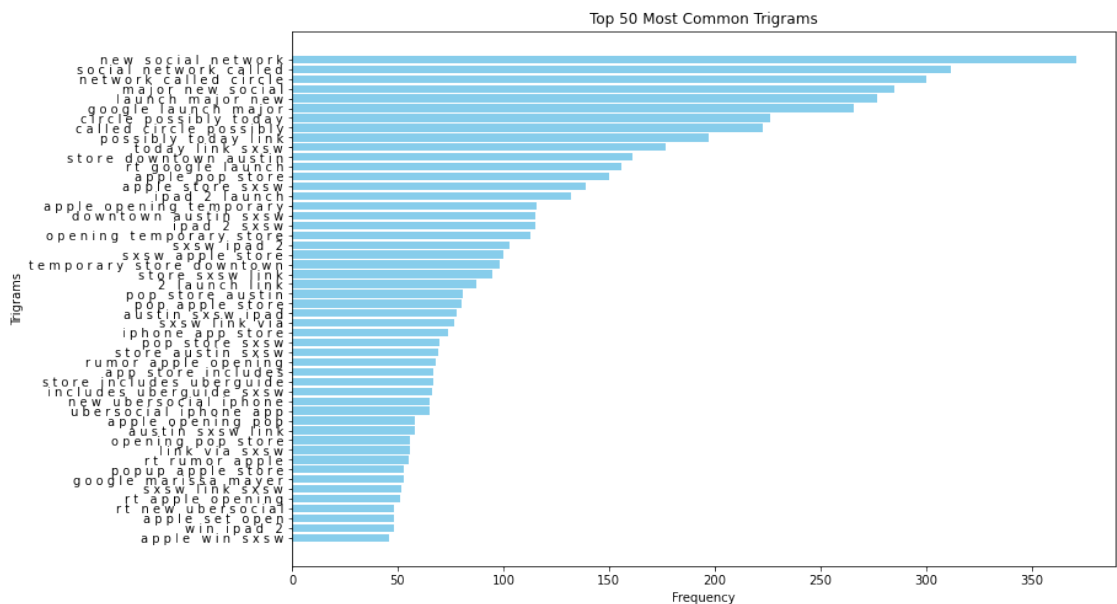
Now let's do the same thing for trigrams.

In [45]:

```

1 # Flatten the list of trigrams
2 all_trigrams = [trigram for sublist in df['trigrams'] for trigram in sublist]
3
4 # Count the frequency of each bigram
5 trigram_freq = Counter(all_trigrams)
6
7 # Get the top 50 most common bigrams
8 top_50_trigrams = trigram_freq.most_common(50)
9
10 # Prepare data for the histogram
11 trigrams, counts = zip(*top_50_trigrams)
12 trigram_labels = [' '.join(trigram) for trigram in trigrams]
13
14 # Plot the histogram
15 plt.figure(figsize=(12, 8))
16 plt.barh(trigram_labels, counts, color='skyblue')
17 plt.xlabel('Frequency')
18 plt.ylabel('Trigrams')
19 plt.title('Top 50 Most Common Trigrams')
20 plt.gca().invert_yaxis() # Invert y-axis to have the highest frequency at the top
21 plt.show()

```




Now I want to see a histogram showing the sentiment associated with each brand.

In [46]:



```
1 plt.figure(figsize=(15, 8))
2 sns.countplot(x='sentiment', hue='brand', data=df)
3 plt.show()
```




We can see most of the tweets are indifferent or positive to the brands and only a relative few are negative. Let's reproduce this graph, but focus on the brand instead of the sentiment.

In [47]:



```
1 sentiment_labels = {1: 'bad', 2: 'neutral', 3: 'good'}
2 df['sentiment_label'] = df['sentiment'].map(sentiment_labels)
3 gradient_palette = sns.color_palette("viridis_r", 3)
4 plt.figure(figsize=(15, 8))
5 sns.countplot(x='brand', hue='sentiment_label', data=df, palette=grad
6 plt.show())
```



Ok, now we can see a little bit more clearly what people are saying about each brand. Most of the tweets about Apple are good and most of the tweets about iPhone, iPad, Google, and Android are neutral. This makes sense, because we could see in the previous chart that most of the tweets are neutral.



We've got some pretty good descriptive and visualizations here. Time to build our models.

First thing we need to do is turn our variables into strings. This will allow our ColumnTransformer class to prepare the text for Term Frequency-Inverse Document Frequency (TF-IDF). This is a stat that determines the specific words importance relative to the all words in the corpus.

```
In [48]: ▶ 1 df['pos_tagged_text'] = df['pos_tagged_text'].astype(str)
2 df['bigrams'] = df['bigrams'].astype(str)
3 df['trigrams'] = df['trigrams'].astype(str)
4 df['tokenized_tweets'] = df['tokenized_tweets'].astype(str)
```

Step 1 for all modeling-defining X and Y variables.

For us, the **'sentiment'** value is going to be the dependant variable and our independant variables are:

- **'pos_tagged_text'**
- **'bigrams'**
- **'trigrams'**
- **'brand_iPhone'**
- **'brand_iPad'**
- **'brand_Apple'**
- **'brand_Google'**
- **'brand_Android'**

```
In [49]: ▶ 1 X = df[['pos_tagged_text', 'bigrams', 'trigrams', 'brand_iPhone', 'br
2 y = df['sentiment']
3
4 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0
```

OK! We're ready to start building some models. Here are the models we are going to build: Let's start with a baseline **Logistic Regression model**. Logistic Regressions are great for modeling probability of outcomes.

We will define preprocessor as the ColumnTransformer function that will use TFIDF vectorization on our text features. This preprocessor will be recycled for every model we build.



```
In [50]: 1 # Step 1: Preprocess the text data
2 preprocessor = ColumnTransformer(
3     transformers=[
4         ('pos_tagged_text', TfidfVectorizer(max_features=5000), 'pos_
5         ('bigrams', TfidfVectorizer(max_features=5000), 'bigrams'),
6         ('trigrams', TfidfVectorizer(max_features=5000), 'trigrams')
7     ]
8 )
9
10 # Step 2: Define and train our Logistic regression model
11 logistic_pipeline = Pipeline(steps=[
12     ('preprocessor', preprocessor),
13     ('classifier', LogisticRegression(max_iter=1000))
14 ])
15
16 logistic_pipeline.fit(X_train, y_train)
17
18 # Step 3: Evaluate the model
19 y_pred = logistic_pipeline.predict(X_test)
20
21 accuracy = accuracy_score(y_test, y_pred)
22 print(f'Accuracy: {accuracy:.2f}')
23 print("Classification Report:")
24 print(classification_report(y_test, y_pred))
```

Accuracy: 0.67

Classification Report:

	precision	recall	f1-score	support
1	0.62	0.18	0.28	117
2	0.69	0.81	0.75	883
3	0.62	0.56	0.59	568
accuracy			0.67	1568
macro avg	0.64	0.51	0.54	1568
weighted avg	0.66	0.67	0.65	1568

Not a bad start. Let's see if we can improve our scores by using **GridSearchCV on the Logistic Regression Model** to find the best parameters. GridSearchCV will run the model with every combination of parameters we give it and then pick the best model. Notice how this code block will also output the best results as well as running the best results and producing the stats. These can take a while to run sometimes.



```
In [51]: ▶ 1 # Define the parameter grid for GridSearchCV
2 param_grid = {
3     'classifier__C': [0.1, 1, 10, 100], # Regularization strength
4     'classifier__solver': ['liblinear', 'saga'] # Solvers
5 }
6
7 # Set up GridSearchCV with the Logistic regression pipeline
8 grid_search = GridSearchCV(estimator=logistic_pipeline, param_grid=pa
9
10 # Train the pipeline with grid search
11 grid_search.fit(X_train, y_train)
12
13 # Get the best parameters and best score
14 best_params = grid_search.best_params_
15 best_score = grid_search.best_score_
16
17 print(f'Best parameters found: {best_params}')
18 print(f'Best cross-validation accuracy: {best_score:.2f}')
19
20 # Make predictions with the best estimator
21 best_logistic_model = grid_search.best_estimator_
22 y_pred_logistic = best_logistic_model.predict(X_test)
23
24 # Evaluate the best model
25 accuracy_logistic = accuracy_score(y_test, y_pred_logistic)
26 print(f'Logistic Regression Accuracy: {accuracy_logistic:.2f}')
27 print("Logistic Regression Classification Report:")
28 print(classification_report(y_test, y_pred_logistic))
```

Best parameters found: {'classifier__C': 1, 'classifier__solver': 'saga'}

Best cross-validation accuracy: 0.66


Logistic Regression Accuracy: 0.67

Logistic Regression Classification Report:

	precision	recall	f1-score	support
1	0.62	0.18	0.28	117
2	0.69	0.81	0.75	883
3	0.62	0.56	0.59	568
accuracy			0.67	1568
macro avg	0.64	0.51	0.54	1568
weighted avg	0.66	0.67	0.65	1568

```
In [52]: 1 # Visualization: Confusion Matrix
2 plt.figure(figsize=(10, 7))
3 ConfusionMatrixDisplay.from_estimator(grid_search, X_test, y_test, cm
4 plt.title('Confusion Matrix - Logistic Regression')
5 plt.show()
```

<Figure size 720x504 with 0 Axes>



Now Let's build a **Decision Tree Classifier** using our preprocessor from above. A Decision Tree is another model great for classifications. Essentially, we create a model that predicts the value of a y-variable by learning simple decision rules inferred from the x-variables.



```
In [53]: ▶ 1 from sklearn.tree import DecisionTreeClassifier
2
3 # Define pipeline for Decision Tree Classifier
4 decision_tree_pipeline = Pipeline(steps=[
5     ('preprocessor', preprocessor),
6     ('classifier', DecisionTreeClassifier())
7 ])
8
9 # Fit the pipeline on the training data
10 decision_tree_pipeline.fit(X_train, y_train)
11
12 # Make predictions on the test data
13 y_pred_decision_tree = decision_tree_pipeline.predict(X_test)
14
15 # Evaluate the model
16 accuracy_decision_tree = accuracy_score(y_test, y_pred_decision_tree)
17 print(f'Decision Tree Classifier Accuracy: {accuracy_decision_tree:.2f}')
18 print("Decision Tree Classifier Classification Report:")
19 print(classification_report(y_test, y_pred_decision_tree))
20
```

Decision Tree Classifier Accuracy: 0.59


Decision Tree Classifier Classification Report:

	precision	recall	f1-score	support
1	0.29	0.16	0.21	117
2	0.65	0.69	0.67	883
3	0.52	0.52	0.52	568
accuracy			0.59	1568
macro avg	0.49	0.46	0.47	1568
weighted avg	0.58	0.59	0.58	1568

Now let's run **GridSearchCV** to optimize our **Decision Tree Classifier**. Notice how this code block will also output the best results as well as running the best results and producing the stats.

Let's take a look at the Decision Tree. This will be overwhelming, but it is still good practice to take a look at it.

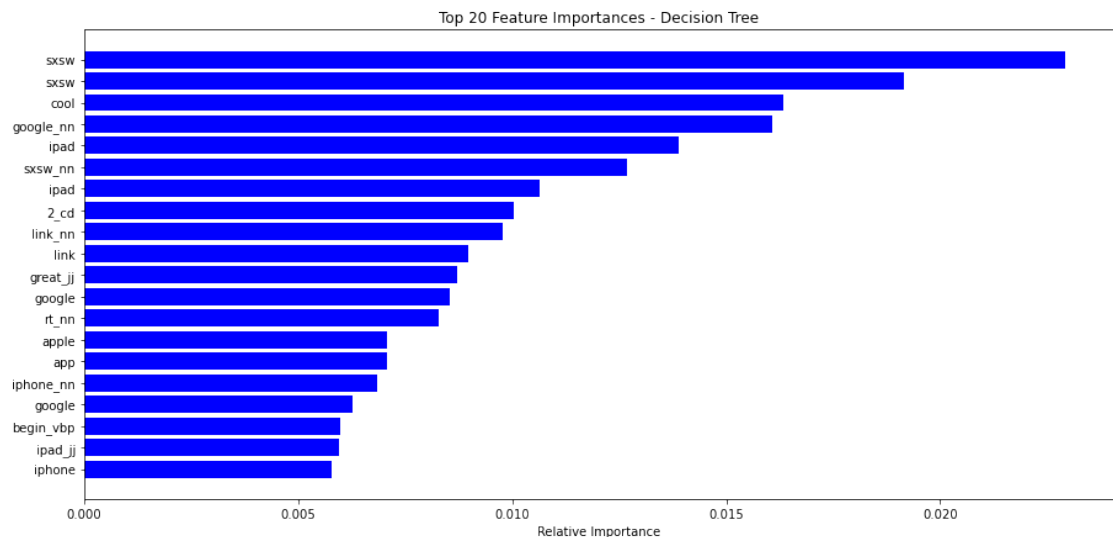

```
In [54]: ▶ 1 from sklearn.tree import plot_tree
2 preprocessor_tree = decision_tree_pipeline.named_steps['preprocessor']
3
4 # Get the feature names
5 feature_names = preprocessor_tree.get_feature_names_out()
6
7 # Plot the decision tree
8 plt.figure(figsize=(20, 10))
9 plot_tree(decision_tree_pipeline.named_steps['classifier'],
10           feature_names=feature_names,
11           class_names=[str(cls) for cls in np.unique(y_train)],
12           filled=True)
13 plt.show()
```



Now we want to view the feature importances in the decision tree. This will help us understand which features had high impact on the model. Notice our brands are all listed, as well as the conference name **SXSW**.



```
In [55]: 1 # Function to extract feature names correctly
2 def get_feature_names(preprocessor):
3     output_features = []
4     for name, transformer, columns in preprocessor.transformers_:
5         if hasattr(transformer, 'get_feature_names_out'):
6             feature_names = transformer.get_feature_names_out()
7             output_features.extend(feature_names)
8         elif hasattr(transformer, 'transformers_'): # Handling nested
9             nested_features = get_feature_names(transformer)
10            output_features.extend(nested_features)
11        else:
12            output_features.extend(columns)
13    return output_features
14
15 # Visualization: Feature Importance using Decision Tree
16 def plot_feature_importances(model, feature_names, top_features=20):
17     importances = model.named_steps['classifier'].feature_importances_
18     indices = np.argsort(importances)[-top_features:]
19
20     plt.figure(figsize=(15, 7))
21     plt.barh(range(len(indices)), importances[indices], color='blue',
22             plt.yticks(range(len(indices)), [feature_names[i] for i in indices]
23             plt.xlabel('Relative Importance')
24             plt.title('Top 20 Feature Importances - Decision Tree')
25             plt.show()
26
27 # Extract feature names from the preprocessor
28 feature_names = get_feature_names(decision_tree_pipeline.named_steps[
29 feature_names = np.array(feature_names)
30
31 # Plot feature importances
32 plot_feature_importances(decision_tree_pipeline, feature_names)
33
```





```
In [70]: 1 # Define the parameter grid for GridSearchCV
2 param_grid = {
3     'classifier__max_depth': [1, 2, 5],
4     'classifier__min_samples_split': [1, 2, 5,],
5     'classifier__min_samples_leaf': [1, 2, 4],
6     'classifier__criterion': ['gini', 'entropy']
7 }
8
9 # Set up GridSearchCV with the decision tree pipeline
10 grid_search = GridSearchCV(estimator=decision_tree_pipeline, param_gr
11
12 # Train the pipeline with grid search
13 grid_search.fit(X_train, y_train)
14
15 # Get the best parameters and best score
16 best_params = grid_search.best_params_
17 best_score = grid_search.best_score_
18
19 print(f'Best parameters found: {best_params}')
20 print(f'Best cross-validation accuracy: {best_score:.2f}')
21
22 # Make predictions with the best estimator
23 best_tree_model = grid_search.best_estimator_
24 y_pred_tree = best_tree_model.predict(X_test)
25
26 # Evaluate the best model
27 accuracy_tree = accuracy_score(y_test, y_pred_tree)
28 print(f'Decision Tree Accuracy: {accuracy_tree:.2f}')
29 print("Decision Tree Classification Report:")
30 print(classification_report(y_test, y_pred_tree))
```

ange [2, int) or a float in the range (0.0, 1.0]. Got 1 instead.

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\byrdw\anaconda3\envs\learn-env\lib\site-packages\sklearn\model_selection\_search.py:979: UserWarning: One or more of the test scores are non-finite: [          nan 0.57623604 0.57623604          nan 0.57623604 0.57623604
          nan 0.57623604 0.57623604          nan 0.5814992 0.58165869
          nan 0.58165869 0.58165869          nan 0.5814992 0.5814992
          nan 0.59473684 0.59473684          nan 0.59425837 0.59441786
          nan 0.59346093 0.59298246          nan 0.56698565 0.56698565
          nan 0.56698565 0.56698565          nan 0.56698565 0.56698565
          nan 0.576874 0.576874          nan 0.576874 0.576874
          nan 0.576874 0.576874          nan 0.58931419 0.58867624
          nan 0.58931419 0.58867624          nan 0.58867624 0.58835726]
warnings.warn(

Best parameters found: {'classifier__criterion': 'gini', 'classifier__max_depth': 5, 'classifier__min_samples_leaf': 1, 'classifier__min_samples_split': 2}
~ . . . . . ~ ~ ~
```

Now we're going to run a **K-Nearest Neighbors Classifier**. KNN is built on the idea that similar data points often have similar features. So during the training phase, KNN calculates the Euclidean Distance between points and groups them by the 'k' number of nearest neighbors. The default is 5 and what we use in our baseline model.



```
In [57]: ▶ 1 from sklearn.neighbors import KNeighborsClassifier
2
3 # Define pipeline for KNN Classifier
4 knn_pipeline = Pipeline(steps=[
5     ('preprocessor', preprocessor),
6     ('classifier', KNeighborsClassifier())
7 ])
8
9 # Fit the pipeline on the training data
10 knn_pipeline.fit(X_train, y_train)
11
12 # Make predictions on the test data
13 y_pred_knn = knn_pipeline.predict(X_test)
14
15 # Evaluate the model
16 accuracy_knn = accuracy_score(y_test, y_pred_knn)
17 print(f'KNN Classifier Accuracy: {accuracy_knn:.2f}')
18 print("KNN Classifier Classification Report:")
19 print(classification_report(y_test, y_pred_knn))
20
```

KNN Classifier Accuracy: 0.63

KNN Classifier Classification Report:

	precision	recall	f1-score	support
1	0.35	0.18	0.24	117
2	0.65	0.83	0.73	883
3	0.61	0.41	0.49	568
accuracy			0.63	1568
macro avg	0.54	0.47	0.48	1568
weighted avg	0.61	0.63	0.60	1568

Now let's run GridSearchCV to optimize our **KNN Classifier**. Notice how this code block will also output the best results as well as running the best results and producing the stats.



```
In [58]: ▶ 1 # Define the parameter grid for GridSearchCV
2 param_grid = {
3     'classifier__n_neighbors': [3, 5, 10],
4     'classifier__weights': ['uniform', 'distance'],
5     'classifier__algorithm': ['auto', 'ball_tree', 'kd_tree'],
6     'classifier__p': [1, 2]
7 }
8
9 # Set up GridSearchCV with the KNN pipeline
10 grid_search_knn = GridSearchCV(estimator=knn_pipeline, param_grid=par
11
12 # Train the pipeline with grid search
13 grid_search_knn.fit(X_train, y_train)
14
15 # Get the best parameters and best score
16 best_params_knn = grid_search_knn.best_params_
17 best_score_knn = grid_search_knn.best_score_
18
19 print(f'Best parameters found for KNN: {best_params_knn}')
20 print(f'Best cross-validation accuracy for KNN: {best_score_knn:.2f}')
21
22 # Make predictions with the best estimator
23 best_knn_model = grid_search_knn.best_estimator_
24 y_pred_knn = best_knn_model.predict(X_test)
25
26 # Evaluate the best model
27 accuracy_knn = accuracy_score(y_test, y_pred_knn)
28 print(f'KNN Accuracy: {accuracy_knn:.2f}')
29 print("KNN Classification Report:")
30 print(classification_report(y_test, y_pred_knn))
```

Best parameters found for KNN: {'classifier__algorithm': 'auto', 'classifier__n_neighbors': 10, 'classifier__p': 2, 'classifier__weights': 'distance'}

Best cross-validation accuracy for KNN: 0.64

KNN Accuracy: 0.64

KNN Classification Report:

	precision	recall	f1-score	support
1	0.46	0.14	0.21	117
2	0.66	0.83	0.73	883
3	0.62	0.46	0.52	568
accuracy			0.64	1568
macro avg	0.58	0.47	0.49	1568
weighted avg	0.63	0.64	0.62	1568

Building Gradient Boosting classifier



```
In [59]: ▶ 1 # Create a pipeline that includes preprocessing and the Gradient Boosting
2 gradient_boosting_pipeline = Pipeline(steps=[
3     ('preprocessor', preprocessor),
4     ('classifier', GradientBoostingClassifier())
5 ])
6
7 # Train the pipeline
8 gradient_boosting_pipeline.fit(X_train, y_train)
9
10 # Make predictions
11 y_pred_gradient_boosting = gradient_boosting_pipeline.predict(X_test)
12
13 # Evaluate the pipeline
14 accuracy_gradient_boosting = accuracy_score(y_test, y_pred_gradient_b
15 print(f'Gradient Boosting Accuracy: {accuracy_gradient_boosting:.2f}')
16 print("Gradient Boosting Classification Report:")
17 print(classification_report(y_test, y_pred_gradient_boosting))
```

Gradient Boosting Accuracy: 0.66


Gradient Boosting Classification Report:

	precision	recall	f1-score	support
1	0.56	0.09	0.15	117
2	0.65	0.93	0.77	883
3	0.73	0.36	0.48	568
accuracy			0.66	1568
macro avg	0.64	0.46	0.47	1568
weighted avg	0.67	0.66	0.62	1568

We can now take a look at the confusion matrix for this model. There are 3 classes, so the model may be slightly more complex than we're used to. We can see that class 2 has so many more instances, that it is the most accurately predicted class.

```
In [60]: 1 # Visualization: Confusion Matrix
2 plt.figure(figsize=(10, 7))
3 ConfusionMatrixDisplay.from_estimator(gradient_boosting_pipeline, X_test)
4 plt.title('Confusion Matrix - Gradient Boosting')
5 plt.show()
```

<Figure size 720x504 with 0 Axes>



Finetuning Gradient Bossting model via GridSearchCV



```
In [61]: ▶ 1 # Define the parameter grid for GridSearchCV
2 param_grid = {
3     'classifier__n_estimators': [5, 10],
4     'classifier__learning_rate': [0.1, 0.2],
5     'classifier__max_depth': [2, 4],
6 }
7
8 # Set up GridSearchCV with the Gradient Boosting pipeline
9 grid_search_gb = GridSearchCV(estimator=gradient_boosting_pipeline, p
10
11 # Train the pipeline with grid search
12 grid_search_gb.fit(X_train, y_train)
13
14 # Get the best parameters and best score
15 best_params_gb = grid_search_gb.best_params_
16 best_score_gb = grid_search_gb.best_score_
17
18 print(f'Best parameters found for Gradient Boosting: {best_params_gb}')
19 print(f'Best cross-validation accuracy for Gradient Boosting: {best_s
20
21 # Make predictions with the best estimator
22 best_gb_model = grid_search_gb.best_estimator_
23 y_pred_gradient_boosting = best_gb_model.predict(X_test)
24
25 # Evaluate the best model
26 accuracy_gradient_boosting = accuracy_score(y_test, y_pred_gradient_b
27 print(f'Gradient Boosting Accuracy: {accuracy_gradient_boosting:.2f}')
28 print("Gradient Boosting Classification Report:")
29 print(classification_report(y_test, y_pred_gradient_boosting))
```

Best parameters found for Gradient Boosting: {'classifier__learning_rate': 0.2, 'classifier__max_depth': 4, 'classifier__n_estimators': 10}

Best cross-validation accuracy for Gradient Boosting: 0.62

Gradient Boosting Accuracy: 0.63

Gradient Boosting Classification Report:

	precision	recall	f1-score	support
1	0.71	0.09	0.15	117
2	0.62	0.97	0.75	883
3	0.78	0.23	0.35	568
accuracy			0.63	1568
macro avg	0.70	0.43	0.42	1568
weighted avg	0.68	0.63	0.56	1568

Now let's take all of those models and use ensemble methods to analyze them even further!
First we will look at the Stacking method.



In [62]:

```

1 # Define base models
2 estimators = [
3     ('decision_tree', DecisionTreeClassifier()),
4     ('knn', KNeighborsClassifier()),
5     ('logistic', LogisticRegression(max_iter=1000)),
6     ('gradient_boosting', GradientBoostingClassifier())
7 ]
8
9 # Create a pipeline that includes preprocessing and the stacking clas
10 stacking_clf = Pipeline(steps=[
11     ('preprocessor', preprocessor),
12     ('classifier', StackingClassifier(
13         estimators=estimators,
14         final_estimator=LogisticRegression(max_iter=1000)
15     ))
16 ])
17
18 # Train the pipeline
19 stacking_clf.fit(X_train, y_train)
20
21 #make predictions
22 y_pred_stacking = stacking_clf.predict(X_test)
23
24 # Evaluate the pipeline
25 stack_accuracy = stacking_clf.score(X_test, y_test)
26 print(f'Stacking Classifier Accuracy: {stack_accuracy:.2f}')
27 print("Stacking Classifier Classification Report:")
28 print(classification_report(y_test, y_pred_stacking))
29

```

Stacking Classifier Accuracy: 0.68

Stacking Classifier Classification Report:

	precision	recall	f1-score	support
1	0.61	0.20	0.30	117
2	0.69	0.85	0.76	883
3	0.68	0.52	0.59	568
accuracy			0.68	1568
macro avg	0.66	0.52	0.55	1568
weighted avg	0.68	0.68	0.67	1568

Now let's look at the Voting method.



```
In [63]: 1 # Similarly, for voting classifier
2 voting_clf = Pipeline(steps=[
3     ('preprocessor', preprocessor),
4     ('classifier', VotingClassifier(
5         estimators=[
6             ('decision_tree', DecisionTreeClassifier()),
7             ('knn', KNeighborsClassifier()),
8             ('logistic', LogisticRegression(max_iter=1000)),
9             ('gradient_boosting', GradientBoostingClassifier())
10        ],
11        voting='soft' # Change to 'hard' for majority voting
12    ))
13 ])
14
15 # Train the pipeline
16 voting_clf.fit(X_train, y_train)
17
18 # Make Predictions
19 y_pred_voting = voting_clf.predict(X_test)
20
21 # Evaluate the pipeline
22 vote_accuracy = voting_clf.score(X_test, y_test)
23 print(f'Voting Classifier Accuracy: {vote_accuracy:.2f}')
24 print("Voting Classifier Classification Report:")
25 print(classification_report(y_test, y_pred_voting))
```

Voting Classifier Accuracy: 0.65

Voting Classifier Classification Report:

	precision	recall	f1-score	support
1	0.52	0.14	0.22	117
2	0.67	0.82	0.74	883
3	0.61	0.49	0.54	568
accuracy			0.65	1568
macro avg	0.60	0.48	0.50	1568
weighted avg	0.64	0.65	0.63	1568

Let's do the same thing, but with with our finetuned models to see the difference. Stacking first.

In [64]:

```

1 # Converting _train and X_test to dataframes for modeling
2 X_train = pd.DataFrame(X_train, columns=['pos_tagged_text', 'bigrams
3 )
4 X_test = pd.DataFrame(X_test, columns=['pos_tagged_text', 'bigrams',
5 )
6
7 # Define the best models from grid search
8 best_tree_model = DecisionTreeClassifier()
9 best_knn_model = KNeighborsClassifier()
10 best_logistic_model = LogisticRegression(max_iter=1000)
11 best_gb_model = GradientBoostingClassifier()
12
13 # Create the stacking classifier pipeline
14 stacking_clf = Pipeline(steps=[
15     ('preprocessor', preprocessor),
16     ('classifier', StackingClassifier(
17         estimators=[
18             ('decision_tree', best_tree_model),
19             ('knn', best_knn_model),
20             ('logistic', best_logistic_model),
21             ('gradient_boosting', best_gb_model)
22         ],
23         final_estimator=LogisticRegression(max_iter=1000)
24     ))
25 ])
26
27 # Train the stacking pipeline
28 stacking_clf.fit(X_train, y_train)
29
30 # Make predictions and evaluate the stacking pipeline
31 y_pred_stacking_ft = stacking_clf.predict(X_test)
32 stack_accuracy = stacking_clf.score(X_test, y_test)
33 print(f'Stacking Classifier Accuracy: {stack_accuracy:.2f}')
34 print("Stacking Classifier Classification Report:")
35 print(classification_report(y_test, y_pred_stacking_ft))
36
37

```

Stacking Classifier Accuracy: 0.69

Stacking Classifier Classification Report:

	precision	recall	f1-score	support
1	0.61	0.20	0.30	117
2	0.69	0.85	0.76	883
3	0.68	0.53	0.59	568
accuracy			0.69	1568
macro avg	0.66	0.53	0.55	1568
weighted avg	0.68	0.69	0.67	1568

Now the voting method.



```
In [65]: ▶ 1 # Create the voting classifier pipeline
2 voting_clf = Pipeline(steps=[
3     ('preprocessor', preprocessor),
4     ('classifier', VotingClassifier(
5         estimators=[
6             ('decision_tree', best_tree_model),
7             ('knn', best_knn_model),
8             ('logistic', best_logistic_model),
9             ('gradient_boosting', best_gb_model)
10        ],
11        voting='soft' # Change to 'hard' for majority voting
12    ))
13 ])
14
15 # Train the voting pipeline
16 voting_clf.fit(X_train, y_train)
17
18 # Make predictions and evaluate the voting pipeline
19 y_pred_voting = voting_clf.predict(X_test)
20 vote_accuracy = voting_clf.score(X_test, y_test)
21 print(f'Voting Classifier Accuracy: {vote_accuracy:.2f}')
22 print("Voting Classifier Classification Report:")
23 print(classification_report(y_test, y_pred_voting))
24
```

Voting Classifier Accuracy: 0.65

Voting Classifier Classification Report:

	precision	recall	f1-score	support
1	0.52	0.11	0.18	117
2	0.67	0.83	0.74	883
3	0.61	0.49	0.54	568
accuracy			0.65	1568
macro avg	0.60	0.48	0.49	1568
weighted avg	0.64	0.65	0.63	1568

Lastly, let's look at Bayesian Statistics and a confusion matrix for these stats. Bayesian stats are great for probabilities. Notice how precise it is for every class!



```
In [66]: ▶ 1 from sklearn.feature_extraction.text import TfidfVectorizer
2 from sklearn.model_selection import train_test_split
3 from sklearn.naive_bayes import MultinomialNB
4 from sklearn.metrics import accuracy_score, classification_report
5
6
7 bayesian_pipeline = Pipeline(steps=[
8     ('preprocessor', preprocessor),
9     ('classifier', MultinomialNB())
10 ])
11
12 bayesian_pipeline.fit(X_train, y_train)
13
14 # Step 4: Make predictions
15 y_pred = bayesian_pipeline.predict(X_test)
16
17 # Step 5: Evaluate the model
18 accuracy = accuracy_score(y_test, y_pred)
19 report = classification_report(y_test, y_pred)
20
21 print(f'Accuracy: {accuracy}')
22 print(f'Classification Report:\n{report}')
```


Accuracy: 0.6594387755102041

Classification Report:

	precision	recall	f1-score	support
1	0.77	0.09	0.15	117
2	0.69	0.81	0.74	883
3	0.60	0.54	0.57	568
accuracy			0.66	1568
macro avg	0.69	0.48	0.49	1568
weighted avg	0.66	0.66	0.64	1568

```
In [67]: 1 import matplotlib.pyplot as plt
2 from sklearn.metrics import ConfusionMatrixDisplay
3
4 # Visualization: Confusion Matrix
5 plt.figure(figsize=(10, 7))
6 ConfusionMatrixDisplay.from_estimator(bayesian_pipeline, X_test, y_te
7 plt.title('Confusion Matrix - Naive Bayes')
8 plt.show()
9
```

<Figure size 720x504 with 0 Axes>



Final Analysis

Our best model was our finetune model that went through Stacking Ensemble method.

Remember:

- Class 1 = Negative
- Class 2= Neutral
- Class 3 = Positive

```
In [68]: 1 print(f'Stacking Classifier Accuracy: {stack_accuracy:.2f}')
2 print("Stacking Classifier Classification Report:")
3 print(classification_report(y_test, y_pred_stacking_ft))
```



Stacking Classifier Accuracy: 0.69

Stacking Classifier Classification Report:

	precision	recall	f1-score	support
1	0.61	0.20	0.30	117
2	0.69	0.85	0.76	883
3	0.68	0.53	0.59	568
accuracy			0.69	1568
macro avg	0.66	0.53	0.55	1568
weighted avg	0.68	0.69	0.67	1568

Here is what we can determine from our results:

- Our model can predict the correct class of tweet(negative, neutral, positive) with 69% accuracy. This is actually fairly accurate as there are 3 classes, so a model that was purely guessing would be accurate roughly 33% of the time.
- Class 1
 - **Precision**- Our model predicts tweets fall into class 1 61% of the time.
 - **Recall**- The instances where the actual value of a tweet is class 1 were correctly identified 20% of the time
 - **F1-score**- Low F1-score is due to the sample size
- Class 2
 - **Precision**- Our model predicts tweets fall into class 2 69% of the time.
 - **Recall**- The instances where the actual value of a tweet is class 2 were correctly identified 85% of the time
 - **F1-score**- High F1-score is due to the sample size
- Class 3
 - **Precision**- Our model predicts tweets fall into class 3 68% of the time.
 - **Recall**- The instances where the actual value of a tweet is class 3 were correctly identified 53% of the time
 - **F1-score**- Moderate F1-score is due to the sample size

Ensemble methods in ML is the process of combining multiple models to create one model that should outperform all others. This works, because the Ensemble method is able to utilize the strengths of the models and is able to sort out the weaknesses of each individual model.

In this notebook, we utilized 2 Ensemble Methods-Stacking and Voting. Here's a brief synopsis of how these work:

Stacking-Multiple models called base learners are trained on the dataset and then the predictions of these models are the input features of a secondary model referred to as a meta-learner.

Voting-the models are all trained independently and then combined via voting. There are two types of voting-hard voting and soft voting. In hard voting, the final prediction is determined by majority vote. In soft voting, each model provides the probability estimate for each class and

the final prediction is made by averaging these probabilities.

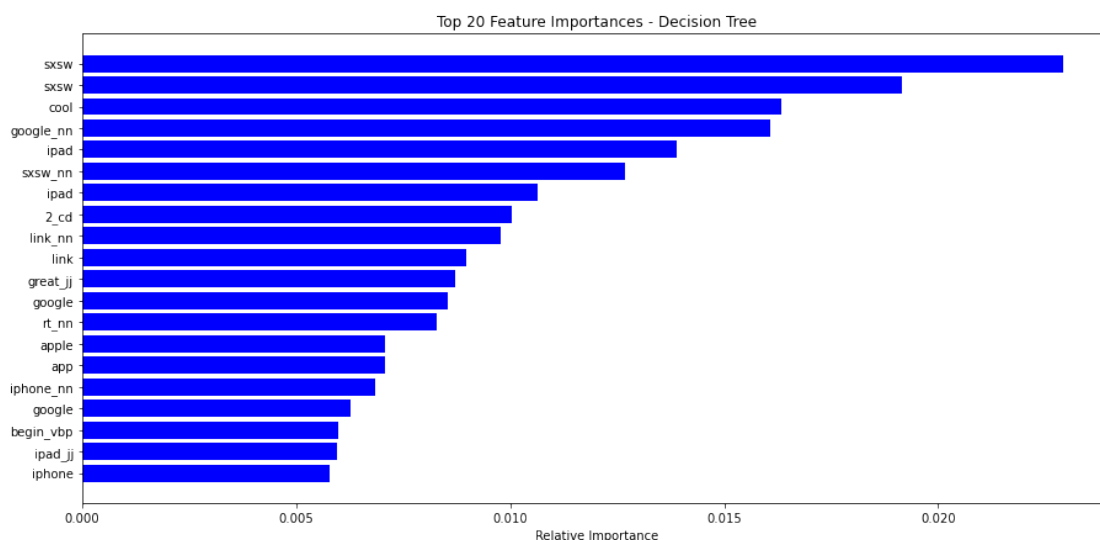


Recomendations

My receomendation to the conference organizers of SXSW would be to allocate as many resources to Apple as they can afford. Clearly, at this time, Apple was the most popular tech brand. iPhone, iPad, Macbook, etc. were all revolutionary products that are still some of the most popular products to this day. Google was the 2nd most popular brand at this conference. Interestingly enough, Google was the most important feature in our modelling, despite it being the 2nd most popular brand.

You can see that all of the Apple products are right up there with Google on feature importance. It's understandable that most of the tweets were nuetral or positive as SXSW is a tech conference. However, it is still imporant to note that Apple and Google were talked about significantly more than Android.

In [69]: `1 plot_feature_importances(decision_tree_pipeline, feature_names)`



Next Steps

We can take more time with our hyperparameter tuning. However, some of these models take so long to run, that it is not an effective us of my local machine.

There are many more sophisticated NLP modelling techniques we can use that than what is in this notebook. For example, we can use Transformersm Recurrent Neural Networks, Transfer Learning, and Contextual Word Embeddings.